

Final Project Submission

Please fill out:

- Student name: Antony Kanai
- Student pace: full time
- Scheduled project review date/time:
- Instructor name: Noah Kandie
- Blog post URL:

Project Overview

My task is to use explorative data analysis to analyze the movies dataset and recommend the course of action to be taken by Microsoft corporation

Objectives

1. Ascertain the most popular genre of movies
2. Establish the most popular publisher
3. Ascertain the most popular rating for movies
4. Ascertain the relationship between production budget and profit realized from the movies

Business Understanding

Microsoft sees all the big companies creating original video content and they want to get in on the fun. They have decided to create a new movie studio, but they don't know anything about creating movies.

Data Understanding

Before, I start the process of understanding my data, I first import the relevant libraries that will enable me to read all my datasets.

These are Pandas, SQLite and numpy

```
In [5]: import pandas as pd
import sqlite3
import numpy as np
import warnings
```

The next step is reading and having a feel of the datasets.

I begin with connecting to the database and viewing the tables in the database

im.db

In [79]:

```
conn = sqlite3.connect('im.db') # Establishing connection to the database
data = pd.read_sql_query('SELECT name from sqlite_master where type= "table";'
data
```

Out[79]:

name

bom.movie_gross.csv

In [4]: bomovies_df = pd.read_csv("zippedData/bom.movie_gross.csv.gz") # reading the data
bomovies_df.head(2) # viewing the first two entries

Out[4]:

	title	studio	domestic_gross	foreign_gross	year
0	Toy Story 3	BV	415000000.0	652000000	2010
1	Alice in Wonderland (2010)	BV	334200000.0	691300000	2010

In []: bomovies_df.info() # getting to know the number of entries and columns and the

rt.movie_info.tsv.gz

In [6]: rt_df = pd.read_csv(("zippedData/rt.movie_info.tsv.gz"), delimiter = "\t") # reading the data
rt_df.head(3) #viewing the first three entries

Out[6]:

	id	synopsis	rating	genre	director	writer	theater_date	dvc
0	1	This gritty, fast-paced, and innovative police...	R	Adventure Classics Drama	William Friedkin	Ernest Tidyman	Oct 9, 1971	S
1	3	New York City, not-too-distant-future: Eric Pa...	R	Drama Science Fiction and Fantasy	David Cronenberg	David Cronenberg Don DeLillo	Aug 17, 2012	
2	5	Illeana Douglas delivers a superb performance ...	R	Drama Musical and Performing Arts	Allison Anders	Allison Anders	Sep 13, 1996	A

In [7]: `rt_df.info()` *#getting to know the number of rows and columns and the datatypes*

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1560 entries, 0 to 1559
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    1560 non-null   int64
1   synopsis              1498 non-null   object
2   rating                1557 non-null   object
3   genre                 1552 non-null   object
4   director              1361 non-null   object
5   writer                1111 non-null   object
6   theater_date          1201 non-null   object
7   dvd_date              1201 non-null   object
8   currency              340 non-null    object
9   box_office            340 non-null    object
10  runtime               1530 non-null   object
11  studio                494 non-null    object
dtypes: int64(1), object(11)
memory usage: 146.4+ KB
```

rt.reviews.tsv.gz

In [12]: *# reading the data and assigning it to rv_df*

```
rv_df = pd.read_csv(("zippedData/rt.reviews.tsv.gz"), delimiter = "\t", encoding='utf-8')
rv_df.tail(3) # viewing the last 3 entries
```

Out[12]:

	id	review	rating	fresh	critic	top_critic	publisher	date
54429	2000	NaN	2/5	rotten	Emanuel Levy	0	EmanuelLevy.Com	July 17, 2005
54430	2000	NaN	2.5/5	rotten	Christopher Null	0	Filmcritic.com	September 7, 2003
54431	2000	NaN	3/5	fresh	Nicolas Lacroix	0	Showbizz.net	November 12, 2002

In [86]: `rv_df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 33988 entries, 0 to 54424
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   id              33988 non-null  int64
1   review          33988 non-null  object
2   fresh           33988 non-null  object
3   critic          33988 non-null  object
4   top_critic      33988 non-null  int64
5   publisher       33988 non-null  object
6   date            33988 non-null  object
7   score           33988 non-null  object
8   outof           27588 non-null  object
dtypes: int64(2), object(7)
memory usage: 2.6+ MB
```

tmdb.movies.csv.gz

In [10]: `tmdb_df = pd.read_csv("zippedData/tmdb.movies.csv.gz") # Reading the tmdb data`
`tmdb_df.head(2) # viewing the first two entries`

Out[10]:

	Unnamed: 0	genre_ids	id	original_language	original_title	popularity	release_date	title
0	0	[12, 14, 10751]	12444	en	Harry Potter and the Deathly Hallows: Part 1	33.533	2010-11-19	Harry Potter and the Deathly Hallows: Part 1
1	1	[14, 12, 16, 10751]	10191	en	How to Train Your Dragon	28.734	2010-03-26	How to Train Your Dragon

In []: `tmdb_df.info()`

tn.movie_budgets.csv.gz

```
In [9]: budgets= pd.read_csv("zippedData/tn.movie_budgets.csv.gz")
        budgets.head(2)
```

```
Out[9]:
```

	id	release_date	movie	production_budget	domestic_gross	worldwide_gross
0	1	Dec 18, 2009	Avatar	\$425,000,000	\$760,507,625	\$2,776,345,279
1	2	May 20, 2011	Pirates of the Caribbean: On Stranger Tides	\$410,600,000	\$241,063,875	\$1,045,663,875

```
In [ ]: budgets.info()
```

Data Preparation

After having a brief overview of the dataset, I started the data preparation process which involved cleaning of the data with the following objectives:

1. Deal with the NaNs/ or missing data
2. Ensure that all columns are in the correct datatype
3. Deal with placeholders if any

In the cells that follow I will conduct data cleaning and ETL for each of the datasets.

bomovies_df

Dealing with missing Values

The first step is to find the proportion of missing values in each of the columns of the bomovies_df

```
In [13]: bomovies_df.isnull().mean()*100
```

```
Out[13]: title          0.000000
         studio         0.147623
         domestic_gross 0.826690
         foreign_gross  39.858282
         year          0.000000
         dtype: float64
```

From the results above, it seems that in the bomovies_df the proportion of missing values for all the columns except the foreign_gross is quite low at less than 1%.

However for foreign gross, the missing value percentage is almost 40%. I assumed that the missing values in foreign_gross means that the movies were sold domestically and did not reach the international market hence their revenue from the international market is 0. Therefore,

I replace the missing values in the foreign_gross column by 0.

I also replaced missing values in domestic_gross column with zero since i assumed that these movies did not sell in the domestic market

The codes below replaces the Nans with 0 in the columns domestic_gross and foreign_gross respectively.

```
In [14]: bomovies_df["domestic_gross"].fillna(0, inplace= True) # replaces all NAns in domestic_gross
bomovies_df["foreign_gross"].fillna(0, inplace= True) # replaces all NAns in foreign_gross
```

```
In [15]: bomovies_df.isna().mean() # Check if the Nulls have disappeared.
```

```
Out[15]: title           0.000000
studio           0.001476
domestic_gross   0.000000
foreign_gross    0.000000
year             0.000000
dtype: float64
```

For the studio column I replaced the missing values with the mode, which is the most occurring studio.

To find the most occurring studio I used the following code:

```
In [16]: bomovies_df["studio"].value_counts().nlargest(1) # Checking for the most common studio
```

```
Out[16]: IFC      166
Name: studio, dtype: int64
```

Since IFC is the most common studio, i replaced the missing values in the studio column with it

```
In [17]: bomovies_df["studio"].fillna("IFC",inplace=True) # replacing missing values in studio
bomovies_df.isna().mean() # checking if the nulls have disappeared
```

```
Out[17]: title           0.0
studio           0.0
domestic_gross   0.0
foreign_gross    0.0
year             0.0
dtype: float64
```

I have now dealt with the missing values in the bomovies_df succesfully.

Converting to the appropriate column datatypes

However, i realized domestic-gross and foreign gross and year are in the wrong datatype and therefore i cast them to the correct datatype

```
In [20]: bomovies_df = bomovies_df = bomovies_df.astype({"domestic_gross": int})
```

```
In [21]: bomovies_df.info() # checking column datatypes have converted succesfull
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3387 entries, 0 to 3386
Data columns (total 5 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   title           3387 non-null   object
 1   studio          3387 non-null   object
 2   domestic_gross  3387 non-null   int32
 3   foreign_gross   3387 non-null   object
 4   year            3387 non-null   int64
dtypes: int32(1), int64(1), object(3)
memory usage: 119.2+ KB
```

rt_df

Dealing with missing values in rt_df

```
In [22]: rt_df.isna().mean()*100 # First find the proportion of missing values in rt_df
```

```
Out[22]: id           0.000000
synopsis    3.974359
rating      0.192308
genre       0.512821
director    12.756410
writer      28.782051
theater_date 23.012821
dvd_date    23.012821
currency    78.205128
box_office  78.205128
runtime     1.923077
studio      68.333333
dtype: float64
```

Dealing with the nulls;

In the dataframe rt_df, i realized that only to columns would be important for my analysis. These are the genre and rating columns. I sliced them from the main dataframe as follows:

```
In [80]: rt_df = rt_df[["genre", "rating"]] # to slice relevant columns from the rt_df
rt_df.head(2)
```

```
Out[80]:
```

	genre	rating
0	Action and Adventure Classics Drama	R
1	Drama Science Fiction and Fantasy	R

I then dealt with the missing values by droppping the rows that had missing values in the column

```
In [24]: rt_df = rt_df.dropna()
```

```
In [25]: rt_df.isna().mean()*100
```

```
Out[25]: genre      0.0  
rating    0.0  
dtype: float64
```

rv_df

Deal with the missing values by dropping all null vlaues in the rows

```
In [26]: rv_df.dropna(inplace =True) # dropping rows with missing values
```

```
In [27]: rv_df.isnull().mean()*100 # checking if the nulls have disappeared.
```

```
Out[27]: id          0.0  
review      0.0  
rating      0.0  
fresh       0.0  
critic      0.0  
top_critic  0.0  
publisher   0.0  
date        0.0  
dtype: float64
```

The ratings column in the rt_reviews (rv_df) can be split into two columns so that we can be able to standardize ratings of all the enries throuough feature engineering. The next cell splits the rating column ito two and displays them as separate columns

```
In [ ]:
```



```
In [28]: new = rv_df["rating"].str.split(pat = "/", n = 1, expand=True) # splitting the
rv_df["score"] = new[0]
rv_df["outof"] = new[1]
rv_df.drop(columns = ["rating"], inplace = True)
rv_df
```

```
Out[28]:
```

	id	review	fresh	critic	top_critic		publisher	date	score	outof
0	3	A distinctly gallows take on contemporary fina...	fresh	PJ Nabarro	0	Patrick Nabarro	November 10, 2018	3	5	
6	3	Quickly grows repetitive and tiresome, meander...	rotten	Eric D. Snider	0	EricDSnider.com	July 17, 2013	C	None	
7	3	Cronenberg is not a director to be daunted by ...	rotten	Matt Kelemen	0	Las Vegas CityLife	April 21, 2013	2	5	
11	3	While not one of Cronenberg's stronger films, ...	fresh	Emanuel Levy	0	EmanuelLevy.Com	February 3, 2013	B-	None	
12	3	Robert Pattinson works mighty hard to make Cos...	rotten	Christian Toto	0	Big Hollywood	January 15, 2013	2	4	
...	
54419	2000	Sleek, shallow, but frequently amusing.	fresh	Gene Seymour	1	Newsday	September 27, 2002	2.5	4	
54420	2000	The spaniel-eyed Jean Reno infuses Hubert with...	fresh	Megan Turner	1	New York Post	September 27, 2002	3	4	
54421	2000	Manages to be somewhat well-acted, not badly a...	rotten	Bob Strauss	0	Los Angeles Daily News	September 27, 2002	1.5	4	
54422	2000	Arguably the best script that Besson has writt...	fresh	Wade Major	0	Boxoffice Magazine	September 27, 2002	3.5	5	
54424	2000	Dawdles and drags when it should pop; it doesn...	rotten	Manohla Dargis	1	Los Angeles Times	September 26, 2002	1.5	5	

33988 rows × 9 columns



In []:

Before I standardize the ratings, i will first convert the column types from string to integer to allow for mathematical computation.

In []:

tmbd_df

Check fo missing values

```
In [81]: tmbd_df.isna().mean()*100 # checking proportion of missing values
```

```
Out[81]: Unnamed: 0      0.0
genre_ids      0.0
id             0.0
original_language  0.0
original_title  0.0
popularity     0.0
release_date   0.0
title          0.0
vote_average   0.0
vote_count     0.0
dtype: float64
```

Perfect! This data has no any missing value.

```
In [30]: tmbd_df["genre_ids"].value_counts() # check for unique values of Genre_ids
```

```
Out[30]: [99]      3700
[]       2479
[18]     2268
[35]     1660
[27]     1145
...
[37, 12]      1
[10752, 878]   1
[28, 53, 10749, 18, 35]  1
[99, 80, 53, 36]  1
[10751, 12, 28]  1
Name: genre_ids, Length: 2477, dtype: int64
```

In []:

budgets_df

Check for missing values

```
In [ ]: budgets.isna().mean()*100 #checks for percentage of missing values.
```

The budgets_df is also complete!

Since the domestic_gross, Production_budget and Worldwide_gross columns are strings, we need to convert them to integer to facilitate feature engineering to make the data more insightful.

```
In [31]: #The codes below convert the columns with the $ sign into integer to facilitate

budgets['domestic_gross'] = budgets['domestic_gross'].apply(lambda x: int(''.join(x.split('$'))))
budgets['production_budget'] = budgets['production_budget'].apply(lambda x: int(''.join(x.split('$'))))
budgets['worldwide_gross'] = budgets['worldwide_gross'].apply(lambda x: int(''.join(x.split('$'))))
```

Perfect! since we have our budget and gross columns as integers, it is possible to create new profit columns. Profit is the difference between production cost and gross revenue

```
In [32]: budgets["Domestic_profit"] = budgets["domestic_gross"] - budgets["production_budget"]
```

```
In [33]: budgets["worldwide_profit"] = budgets["worldwide_gross"] - budgets["production_budget"]
budgets.info() # check if we have the correct datatypes
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5782 entries, 0 to 5781
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    5782 non-null   int64
1   release_date          5782 non-null   object
2   movie                 5782 non-null   object
3   production_budget     5782 non-null   int64
4   domestic_gross        5782 non-null   int64
5   worldwide_gross       5782 non-null   int64
6   Domestic_profit       5782 non-null   int64
7   worldwide_profit      5782 non-null   int64
dtypes: int64(6), object(2)
memory usage: 361.5+ KB
```

Data Analysis

After prepaing the data and making sure that it is not dirty, I delved into data analysis. In thi section i will try to make sense of the data. I will ty to merge datasets to come up with more insightful analysis and to create a story that microsoft would definitely buy in.

My Data Analysis will focus on Establishing the following:

1. Which is the most popular genre of movies
2. Which is the most popular studio
3. Which rating is most preferred
4. wh

Most popular Genre of Movies

From the rt_movies dataset I can establish the top 5 genres of movies

```
In [34]: rt_df.genre.value_counts().nlargest(5)
```

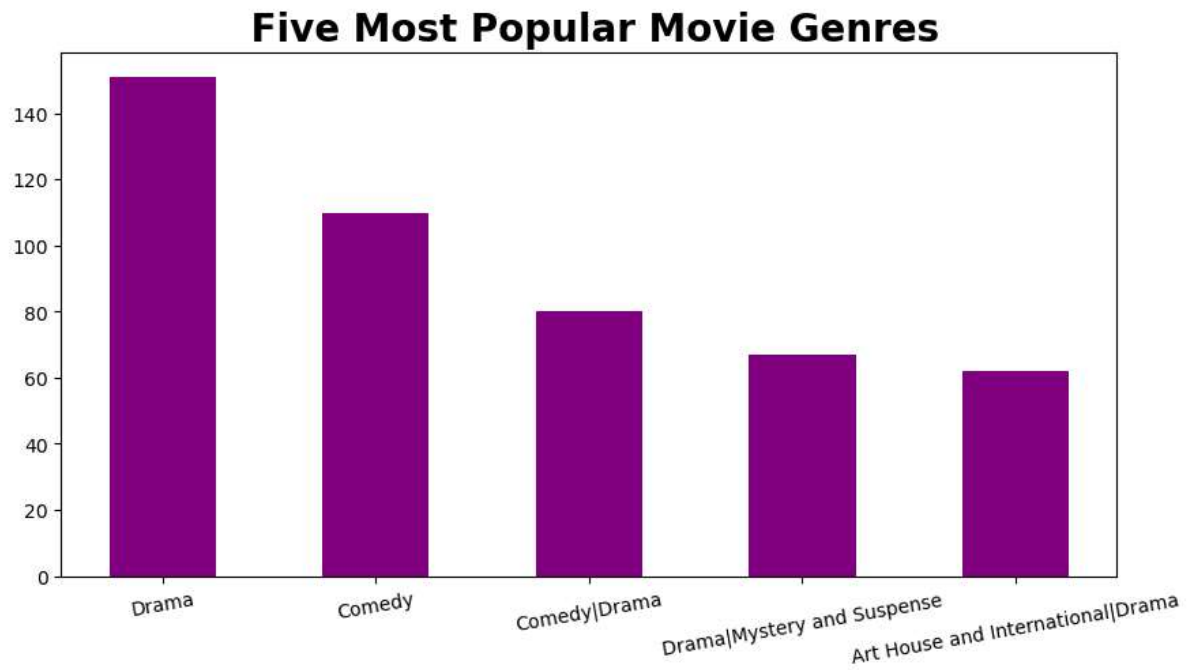
```
Out[34]: Drama                151
Comedy                110
Comedy|Drama             80
Drama|Mystery and Suspense  67
Art House and International|Drama  62
Name: genre, dtype: int64
```

From the above code, the most popular genre of movies is Drama, followed by Comedy, then comedy|Drama and so on... This can be better presetnteed in the visualization below:

```
In [35]: import matplotlib.pyplot as plt # Importing the library neccesaty to create vi.
%matplotlib inline
import seaborn as sns
```

```
In [36]: # Getting the X and Y values
x = rt_df["genre"].value_counts().head().index.tolist()
y = list(rt_df["genre"].value_counts().nlargest(5))
```

```
In [37]: # Ploting Most popular Genres and their frequencies
fig,ax = plt.subplots(figsize = (10,5))
plt.bar(x, y, color = "Purple", width = 0.5)
plt.xticks(rotation = 10);
plt.title(" Five Most Popular Movie Genres", fontsize = 20, fontweight = "bold")
```



Most Popular studio

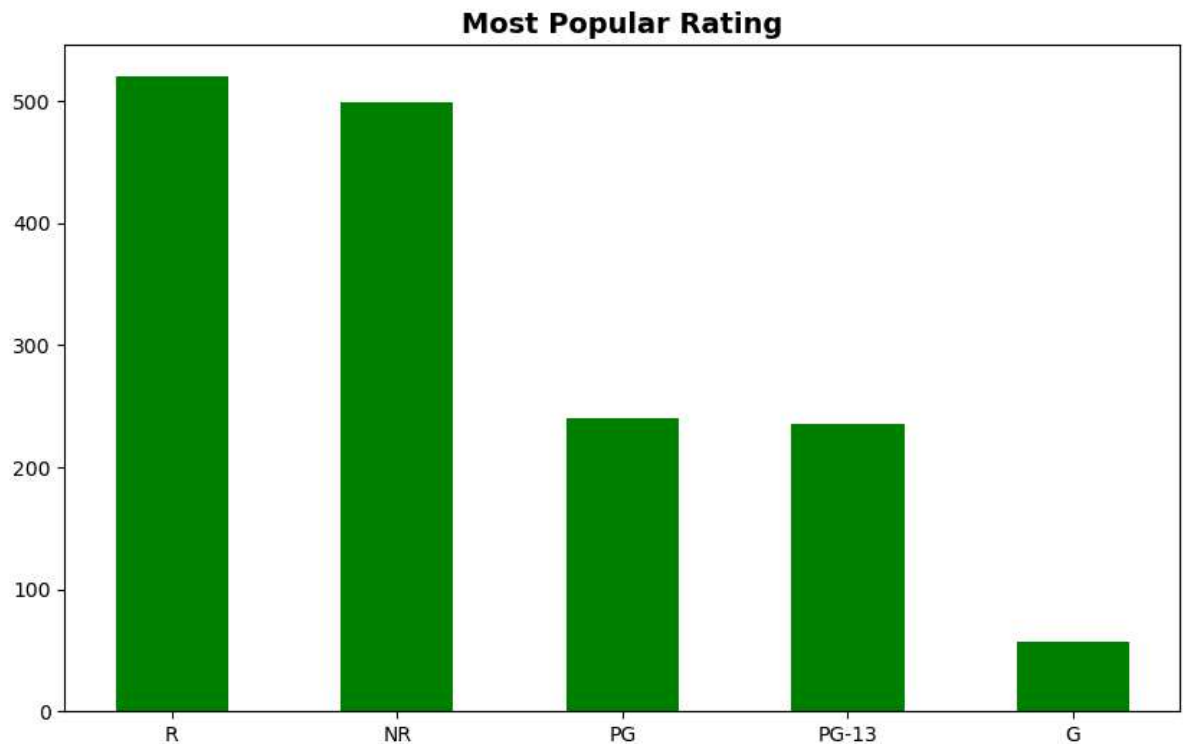
I can be able to find the most popular studio from the bomovies_df

```
In [ ]: rt_df["rating"].value_counts()
```

These can be visually presented as:

```
In [38]: rating= rt_df["rating"].value_counts().head().index.tolist()
frequency=list(rt_df["rating"].value_counts().nlargest(5))
```

```
In [39]: fig, ax = plt.subplots(figsize= (10,6))  
  
plt.bar(rating, frequency, color = "green", width =0.5); # plotting ratings and  
  
plt.title("Most Popular Rating", fontsize = 14, fontweight = "bold");
```



Most Popular Rating

Is There a Relationship Between Movie Production Budget and Profits Realised

In order to ascertain whether production budget affects profitability, I calculated correlation between production budget and profit realised domestically and worldwide

In [40]: `budgets.corr()`

C:\Users\user\AppData\Local\Temp\ipykernel_10264\2354375143.py:1: FutureWarning: The default value of `numeric_only` in `DataFrame.corr` is deprecated. In a future version, it will default to `False`. Select only valid columns or specify the value of `numeric_only` to silence this warning.


`budgets.corr()`

Out[40]:

	id	production_budget	domestic_gross	worldwide_gross	Domestic_profit
id	1.000000	-0.035278	0.008255	-0.009422	0.040832
production_budget	-0.035278	1.000000	0.685682	0.748306	0.099742
domestic_gross	0.008255	0.685682	1.000000	0.938853	0.792663
worldwide_gross	-0.009422	0.748306	0.938853	1.000000	0.656626
Domestic_profit	0.040832	0.099742	0.792663	0.656626	1.000000
worldwide_profit	-0.001172	0.608752	0.926605	0.981811	0.756626

In the correlation matrix above, The correlation between Production budget and worldwide_profit is positive and higher than that between production budget and domestic profit. Below is scatter plot showing the relationship

In [82]: `sns.set_style("darkgrid")`
`fig, axes = plt.subplots(1,2,figsize=(15,5))`
`sns.scatterplot(x= budgets["production_budget"], y =budgets["Domestic_profit"])`
`sns.scatterplot(x= budgets["production_budget"], y =budgets["worldwide_profit"])`



```
In [ ]: sns.set_style("darkgrid")
fig, axes = plt.subplots(1,2,figsize= (15,5))
sns.scatterplot(x= budgets["production_budget"], y =budgets["domestic_gross"],
sns.scatterplot(x= budgets["production_budget"], y =budgets["worldwide_gross"]
```

Most Popular Genre_ids

```
In [84]: # grouping by genre-ids and aggregating by max()

grouped1= tmdb_df.groupby(["genre_ids"]).max().sort_values(by="popularity", as
grouped1.head()
```

Out[84]:

	Unnamed: 0	id	original_language	original_title	popularity	release_date	ti
genre_ids							
[12, 28, 14]	24034	522417	zh	奇门遁甲	80.773	2018-10-23	Thousa Faces Dur
[28, 53]	26399	569869	th	우는 남자	78.123	2018-12-20	Your Mc
[28, 12, 16, 878, 35]	23812	324857	en	Spider-Man: Into the Spider-Verse	60.534	2018-12-14	Spider-M: Into 1 Spid Ver
[28, 12, 14]	24318	525135	zh	西游记之孙悟空三打白骨精	53.783	2018-12-21	Warcr
[878, 28, 12]	24924	521323	en	Wastelanders	50.289	2018-02-02	Wastelanc

Popularity By Language

```
In [85]: # Grouping by Original_Language
grouped = tmdb_df.groupby(["original_language"]).sum().sort_values(by="popularity")
grouped.head()
```

C:\Users\user\AppData\Local\Temp\ipykernel_10264\4174702788.py:2: FutureWarning: The default value of numeric_only in DataFrameGroupBy.sum is deprecated. In a future version, numeric_only will default to False. Either specify numeric_only or select only columns which should be valid for the function.

```
grouped = tmdb_df.groupby(["original_language"]).sum().sort_values(by="popularity", ascending=False)
```

```
Out[85]:
```

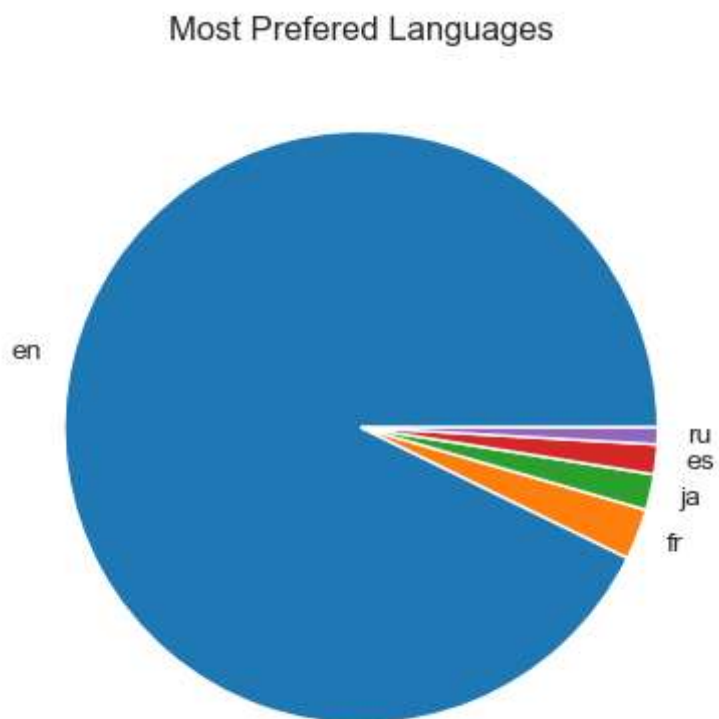
	Unnamed: 0	id	popularity	vote_average	vote_count
original_language					
en	312028215	7005029780	71895.155	138662.0	4874990
fr	5744495	118048030	2155.574	3130.8	75337
ja	3769256	70813222	1513.434	1809.1	54774
es	6070196	127264882	1257.725	2874.3	29396
ru	2859417	64494601	708.220	1579.4	4901

```
In [63]: language = ["en", "fr", "ja", "es", "ru"]
Popularity= [71896, 2155, 1513, 1257, 708]
```

```
In [60]: grouped["popularity"].head()
```

```
Out[60]: original_language
en      71895.155
fr      2155.574
ja      1513.434
es      1257.725
ru       708.220
Name: popularity, dtype: float64
```

```
In [72]: fig,ax = plt.subplots()
plt.pie(Popularity, labels = language)
plt.title("Most Preferred Languages");
```



Merging Datasets


```
In [83]: Merged_df= pd.merge(rv_df, budgets) #Merging datasets with common columns
Merged_df.head()
```

Out[83]:

	id	review	fresh	critic	top_critic	publisher	date	score	outof	release_date
0	3	A distinctly gallows take on contemporary fina...	fresh	PJ Nabarro	0	Patrick Nabarro	November 10, 2018	3	5	Jun 7, 2019
1	3	A distinctly gallows take on contemporary fina...	fresh	PJ Nabarro	0	Patrick Nabarro	November 10, 2018	3	5	Nov 21, 2018
2	3	A distinctly gallows take on contemporary fina...	fresh	PJ Nabarro	0	Patrick Nabarro	November 10, 2018	3	5	Apr 8, 2005
3	3	A distinctly gallows take on contemporary fina...	fresh	PJ Nabarro	0	Patrick Nabarro	November 10, 2018	3	5	Oct 5, 2018
4	3	A distinctly gallows take on contemporary fina...	fresh	PJ Nabarro	0	Patrick Nabarro	November 10, 2018	3	5	Feb 18, 2005

```
In [ ]:
```

```
In [77]: publisher= Merged_df["publisher"].value_counts().head().index.tolist()
frequency=list(Merged_df["publisher"].value_counts().nlargest(5))
sns.barplot(x=publisher, y = frequency, data = Merged_df, palette='OrRd')
plt.xticks(rotation = 30);
plt.title("TOP PUBLISHERS");
```



Summary

The Visualizations above inform how microsoft will begin their movie-Production journey. With the insights we can be sure that Microsoft wil get it right in data analysis

In []: